

# Resource Competition in Multi-Agent Reinforcement Learning:

A Literature Review

---

Anastasia Chernavskaia, Moritz Peist, Nicolas Rauth

June 29, 2025



Barcelona School of Economics

Introduction to Multi-Agent Reinforcement Learning

Adversarial Environments

PPO Algorithm in MARL

Environment Preprocessing

Training and Evaluation

Implementation Highlights

Challenges and Solutions

References

# Introduction to Multi-Agent Reinforcement Learning

---

- Extension of single-agent RL to environments with **multiple learning agents**
- Agents must learn to interact, coordinate, or compete with each other
- Each agent's optimal policy depends on other agents' policies
- Introduces challenges: non-stationarity, coordination, credit assignment

## Key Characteristics:

- Simultaneous learning and adaptation
- Emergent behaviors from agent interactions
- Applications: robotics, autonomous vehicles, game theory, economics

# Adversarial Environments

---

## Environment Setup:

- **Agents:** 1 adversary + 1 cooperative agent
- **Landmarks:** Fixed black landmarks in the environment
- **Objective:** Adversary tries to reach landmarks, agent tries to prevent it

## Competitive Dynamics:

- Zero-sum game structure
- Adversary reward =  $-$  Agent reward
- Tests robustness of learned policies
- Requires strategic thinking and adaptation

# PPO Algorithm in MARL

---

## Why PPO for MARL?

- **Stability:** Prevents large policy updates that could destabilize training
- **Sample Efficiency:** Reuses experience through multiple epochs
- **Simplicity:** Easy to implement and tune compared to TRPO

## PPO Objective Function:

$$L^{CLIP}(\theta) = \mathbb{E}_t \left[ \min \left( r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t \right) \right]$$

where  $r_t(\theta) = \frac{\pi_\theta(a_t|s_t)}{\pi_{\theta_{old}}(a_t|s_t)}$  and  $\epsilon = 0.2$



## Network Architecture:

- **Hidden layers:** [256, 256, 128] neurons
- **Activation:** Hyperbolic tangent (Tanh)
- **Policy:** Actor-Critic with shared features

## Training Parameters:

- **Learning rate:**  $3 \times 10^{-4}$
- **Batch size:** 512, **Steps per update:** 256
- **Discount factor:**  $\gamma = 0.98$
- **GAE lambda:**  $\lambda = 0.95$
- **Entropy coefficient:** 0.01 (encourages exploration)

# Environment Preprocessing

---

## Environment Pipeline:

1. **PettingZoo MPE:** Multi-Particle Environment framework
2. **Black Death:** Removes dead agents from observation/action spaces
3. **Pad Observations:** Ensures consistent observation dimensions
4. **Flatten:** Converts structured observations to flat vectors
5. **AEC to Parallel:** Transforms turn-based to simultaneous actions
6. **Vectorization:** Enables batch processing for efficiency

## Benefits:

- Standardized interface for RL algorithms
- Improved computational efficiency

# Training and Evaluation

---

## Training Configuration:

- **Total timesteps:** 500,000
- **Environment:** Continuous action space
- **Monitoring:** Weights & Biases (wandb) integration
- **Logging:** TensorBoard for real-time metrics

## Key Metrics Tracked:

- Episode rewards (individual and collective)
- Policy and value function losses
- Episode lengths and success rates
- Agent coordination measures

## Evaluation Methods:

- **Animated GIFs:** Show agent behavior evolution
- **Training dashboards:** Real-time metric visualization
- **Performance plots:** Reward trends and learning curves

## Behavioral Analysis:

- Agent positioning strategies
- Adversarial adaptation patterns
- Coordination emergence over time
- Robustness to opponent strategies

## Implementation Highlights

---

## Key Libraries and Frameworks:

- **Stable-Baselines3:** PPO implementation
- **PettingZoo:** Multi-agent environment suite
- **SuperSuit:** Environment preprocessing wrappers
- **PyTorch:** Neural network backend
- **Weights & Biases:** Experiment tracking

## Code Structure:

- Modular design for reusability
- Custom visualization library
- Reproducible experimental setup
- Comprehensive logging and monitoring



# Challenges and Solutions

---

## Non-stationarity:

- **Problem:** Environment changes as other agents learn
- **Solution:** PPO's stable policy updates + experience replay

## Credit Assignment:

- **Problem:** Which agent caused the reward?
- **Solution:** Individual rewards + shared value function

## Scalability:

- **Problem:** Exponential growth in joint action space
- **Solution:** Vectorized environments + parallel training

# References

---

